



IROS2012 Vila Moura, Algarve, Portugal

## PCL :: Segmentation - Planes, Clusters & More

Alexander J B Trevor, Georgia Institute of Technology

October 13, 2012


# Motivation

- Segmentation is used for many applications
  - Object Detection & Recognition
  - Mapping
  - Obstacle avoidance
- Make a fast, general approach focused on Kinect data

# My application: What should maps look like?

- Let's focus on fetch & carry tasks
- Need to know task-relevant features:
  - Structures: walls, tables, shelves
  - Objects: things we want to manipulate or interact with
  - People: users (not addressed here)






A. J. B. Trevor, J. G. Rogers III, C. Nieto-Granda, H. I. Christensen, "Planar Surface SLAM with 3D and 2D Sensors", ICRA, 2012.

# Limitations of the approach

- RANSAC takes a lot of time, which is okay for scanning lasers, but not for Kinect
- Convex hulls poorly represent the shape of many surfaces
- We can do better if we do smarter processing!





# How can we improve this?


- The previous approach was fast enough for laser data. What if we want to use a Kinect on a robot, or even handheld sensor?
- Convex hulls don't represent shapes well, and alpha shapes can be problematic (choosing alpha)
- So, let's not throw away organized structure, and do smarter processing

# Fast Segmentation of Organized Point Cloud Data

- First step: calculate surface normals
- Use Integral Image Normal Estimation
- Takes about 70 ms per frame (covariance matrix method), or 30-40 ms (simple gradient method)

# Connected Components


- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc






# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc




# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc




# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc




# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc




# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc



# Connected Components

- Next, do connected component labeling
- Implementation allows various different comparisons between neighboring pixels
- Different comparison functions allow different types of segmentation: planes, objects, color blobs, etc




# Multi Plane Segmentation

- Comparison function for plane segmentation:

$$\begin{aligned} n_1 \cdot n_2 &< \theta_{thresh} \\ |d_1 - d_2| &< d_{thresh} \end{aligned}$$

- Fit planes to regions with  $> \text{min\_inliers}$






# Plane Coefficient Comparator

```
/** \brief Compare points at two indices by their plane equations. True if the angle between the normals is less than the angular threshold,
 * and the difference between the d component of the normals is less than distance threshold, else false
 * \param idx1 The first index for the comparison
 * \param idx2 The second index for the comparison
 */
virtual bool
compare (int idx1, int idx2) const
{
    float threshold = distance_threshold_;
    if (depth_dependent_)
    {
        Eigen::Vector3f vec = input_->points[idx1].getVector3fMap ();

        float z = vec.dot (z_axis_);
        threshold *= z * z;
    }
    return ( (fabs ((*plane_coeff_d_)[idx1] - (*plane_coeff_d_)[idx2]) < threshold)
            && (normals_->points[idx1].getNormalVector3fMap ().dot (normals_->points[idx2].getNormalVector3fMap () ) > angular_threshold_ ) );
}
```


# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec



# Boundary Detection


- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec






# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Boundary Detection

- Organized data makes it easy to find the boundary of segmented regions
- Trace outer contour of a region
- Example convex hull timing (80k pts): 0.012 sec
- Example boundary timing (80k pts): 0.000048 sec




# Refinement of planar regions

- Normals near surface edges can be noisy
- We can address this issue with an additional two passes
- Instead of comparing neighboring pixels values, we compare to the planar model of neighboring pixels




# Planar Refinement

- Given a pixel that is an inlier to a planar region, we can extend the region to adjacent points using point to plane distance
- Requires two additional passes



# Planar Refinement

- Given a pixel that is an inlier to a planar region, we can extend the region to adjacent points using point to plane distance
- Requires two additional passes



```

/** \brief Compare two neighboring points
 * \param[in] idx1 The index of the first point.
 * \param[in] idx2 The index of the second point.
 */
virtual bool
compare (int idx1, int idx2) const
{
    int current_label = labels_->points[idx1].label;
    int next_label = labels_->points[idx2].label;

    if (!((*refine_labels_)[current_label] && !(*refine_labels_)[next_label]))
        return (false);

    const pcl::ModelCoefficients& model_coeff = (*models_)[(*label_to_model_)[current_label]];

    PointT pt = input_->points[idx2];
    double ptp_dist = fabs (model_coeff.values[0] * pt.x +
                             model_coeff.values[1] * pt.y +
                             model_coeff.values[2] * pt.z +
                             model_coeff.values[3]);

    // depth dependent
    float threshold = distance_threshold_;
    if (depth_dependent_)
    {
        Eigen::Vector3f vec = input_->points[idx1].getVector3fMap ();

        float z = vec.dot (z_axis_);
        threshold *= z * z;
    }


    return (ptp_dist < threshold);
}

```




# Using Color Information

- We can easily segment on color and geometric information at little additional cost
- We can use the same comparison as we did for planes, plus an additional constraint that neighboring pixels have similar color




# Euclidean Cluster Extraction

- Use planar regions extracted in the first pass as a mask
- Simply compare euclidean distance between neighboring points



# Using Boundaries and Holes


- Since we have organized data, we can also notice missing / NAN data
- This can be useful for detecting small shiny objects which may not give good depth returns





# Robust to Lighting Conditions

- Comparison functions that don't use color data work fine in the dark!




# Timing Improvements

- Fair comparison is not so easy due to the large number of parameters.
- For this comparison, I chose to use Iterative RANSAC as used in my previous mapping work, 1000 iterations
- Clusters RANSAC planes to separate coplanar disjoint regions (e.g. multiple tables)
- Computes convex hull, since the MPS gives a boundary
- Data set is kinect data from the “yesterday’s sushi” area, 508 frames
- Timings are reported for plane extraction + hull / boundary computation only, so this is on top of normal estimation


# Timing Improvements

- MultiPlaneSegmentation, PlaneCoefficientComparator, SegmentAndRefine:
  - mean = 0.0332 seconds, stddev = 0.0023
- RANSAC (1000 iterations, dominant plane, 2.5cm voxelized cloud):
  - mean = 0.32 seconds, stddev = 0.34
- RANSAC (1000 iterations, all planes in the scene, 2.5cm voxelized cloud):
  - mean = 1.52 seconds, stddev = 1.79
- RANSAC (1000 iterations, all planes in the scene, 1cm voxelized cloud):
  - mean = 6.18 seconds, stddev = 6.62







Plane Segmentation Timing Comparison



Plane Segmentation Timing Comparison





# Code Example

```
// Segment planes
pcl::OrganizedMultiPlaneSegmentation<PointT, pcl::Normal, pcl::Label> mps;
mps.setMinInliers (10000);
mps.setAngularThreshold (0.017453 * 2.0); // 2 degrees
mps.setDistanceThreshold (0.02); // 2cm
mps.setInputNormals (normal_cloud);
mps.setInputCloud (cloud);
std::vector<pcl::PlanarRegion<PointT> > regions;
mps.segmentAndRefine (regions);

for (size_t i = 0; i < regions.size (); i++)
{
    Eigen::Vector3f centroid = regions[i].getCentroid ();
    Eigen::Vector4f model = regions[i].getCoefficients ();
    pcl::PointCloud<PointT> boundary_cloud;
    boundary_cloud.points = regions[i].getContour ();
    printf ("Centroid: (%f, %f, %f)\n Coefficients: (%f, %f, %f, %f)\n Inliers: %d\n",
           centroid[0], centroid[1], centroid[2],
           model[0], model[1], model[2], model[3],
           boundary_cloud.points.size ());
}
```

- See PCL's Organized Segmentation Demo for more:  
[pcl\\_organized\\_segmentation\\_demo](#)

# Polygon Merge

- Given two 3D polygons (not necessarily exactly co-planar) and plane equations, compute the new polygon and merged plane equation
- Project to 2D
- Use boost::geometry's fast 2D polygon boolean operations
- Coming soon to PCL

# Mapping Planar Surfaces from multiple views: Planar Polygon Fusion

- Considering multiple frames is more interesting than single views
- Given a set of frames and sensor locations (from odom, AMCL, etc) we can fuse planar surfaces across views

